

[\[Return to Main Page\]](#) Investigating Interrupts by Garth Wilson
[\[Up to Tutorials and Primers\]](#)

Table of Contents

- 1 Learning about Interrupts
 - [1.1](#) Introduction
 - [1.2](#) Scenarios Where Interrupts are Useful
 - [1.3](#) Getting to the Nuts and Bolts
 - [1.4](#) Setting Up an Interrupt
- 2 Special Types of Interrupts
 - [2.1](#) NMI: Non-Maskable Interrupt
 - [2.2](#) BRK: Software Interrupt
 - [2.3](#) WAI: Faster Interrupt Service on the 65C02
- 3 Sample Applications
 - [3.1](#) Interrupt Support for RS-232 Receive
 - [3.2](#) Interrupt-Driven Arbitrary Waveform Generator
 - [3.3](#) Simple Analog Output from the 6522's Serial Port

[Appendix A: Interrupt Performance](#)

1.1 INTRODUCTION

Forum posts indicate that many members are uncomfortable with interrupts or even misunderstand why they are used. This primer was written to help with basic interrupt understanding and how to implement interrupts on 6502 systems, without getting into too much detail of individual hardware setups.

For further reading, see the Synertek programming manual [on 6502.org](http://6502.org) or at <http://users.telenet.be/kim1-6502/6502/proman.html>, pages 129-146. These two locations have the same thing in different forms. Then there's [WDC's outstanding programming manual](#) which I can't pass up an opportunity to recommend. Virtually all programming manuals will give some treatment to interrupts. What I have tried to additionally provide here is:

- more examples of what interrupts really are (and are not)
- tell why we use them, in layman's terms
- tell why certain things are done the way they are
- explain some of the insides that must be considered
- give some real-life assembly-language examples of setting up, prioritizing, and servicing interrupts
- give an idea of how the 6502's interrupt performance compares to that of other processors.

Interrupt-Service Routines (ISRs, also called "interrupt handlers") can range from very simple to very complex. Many things can be done with hardware prioritizing of interrupts, interrupt handling in high-level languages, and other things that have not been mentioned here. This is only intended to be a primer, not an exhaustive resource. For the beginner, it could become very difficult to follow more of these things without getting away from the subject of interrupts and giving a lot of attention to other issues of hardware and systems that have little relevance to the reader's application.

If you have more questions about interrupts, or if you find anything questionable, inaccurate, or confusing in this primer, please contact me. I can usually get back to you within a few hours. I will be happy to help further if I can, and I definitely would want to correct any errors anyone might find here.

Garth Wilson
wilsonmines@dslextrreme.com

1.2 SCENARIOS WHERE INTERRUPTS ARE USEFUL

Consider some scenarios:

Scenario A: Suppose you have data bytes coming to the computer via a serial port, possibly through a 6551 ACIA. Further, suppose that servicing the port does not take as much time as there is between bytes coming in. Further, suppose that there are occasional pauses of unpredictable duration in the data.

You could have a routine that does nothing but continually check the serial port as long as it's possible that data could be arriving and the port could need servicing. This is called "polling" the port. Of course that way, the computer wouldn't get

anything else done. It's tied up until the program decides to stop accepting data.

Here's a slightly more productive way to stay on top of the job of reading the data from the serial port and getting the port ready to accept the next byte. You can have a program running that does something else useful but keeps checking back with the port (ie, polling it) periodically to see if the next byte has finished coming in yet. This would probably be done in a subroutine that gets called at different places in the main program. If a byte has finished coming in, the subroutine branches to where it services the port. When the subroutine ends, execution resumes in the main program at the instruction following the subroutine call (ie, the instruction following the JSR). This method still does not make the best use of processor time, and requires that the main program have the JSR's to poll the port sprinkled possibly throughout much of the program.

The much better way to handle the situation uses interrupts. Set the serial port up to tell the processor when it needs attention. Then the background program doesn't need to be written with anything special to poll the port, and the port will get serviced without having to wait for the background program to get to one of the subroutine calls to poll the port. (An actual coding example of this will be given near the end.)

Scenario B: Suppose your computer needs to do a particular job every so many milliseconds or microseconds, and it is rather critical that the timing be kept accurate.

You could have a routine with a timing delay loop that takes a desired, predictable amount of time to finish and drop through so the routine can do the job before starting back into the timing loop again. But now suppose that depending on various conditions, this regularly scheduled job doesn't always take the same amount of time, so the routine has to give differing parameters to the delay loop so its next delay is correct. Adding up the clock cycles of the various instructions involved gets complicated quickly, especially with all the conditional branches, and your resulting timing will only be somewhat close at best. After all your hard work, you still won't have nearly the accuracy that crystal control should give you.

The much better way to handle the situation uses interrupts. Use an IC with a timer implemented in hardware. Set it up to roll over every so many milliseconds or microseconds as the job requires, and to tell the processor when it rolls over. The timer's roll-over rate will remain at whatever value your program gave it, regardless of what the processor is doing. That means that the interrupts will be perfectly spaced and timed regardless of how long it takes the processor to do what it's supposed to do each time the timer rolls over (assuming the interrupt-service job is done before the next roll-over).

A benefit of using the interrupts is that now it won't matter what other job the processor is doing while the timer is running. The processor can be busy doing something else productive until the timer "blows the whistle," so to speak. By the way, if the timer's hardware limitations make it unable to go long enough between roll-overs, you can set your interrupt-service routine (ISR) to count roll-overs and then do the job every so many roll-overs. (The first coding example later will deal with a similar case, although the job will only be to increment a two-byte variable, using it only to mark elapsed time for other routines to refer to.)


Scenario C: Now suppose you have both scenario A and scenario B at the same time.

You can see that running A, when you don't know exactly how fast the data will be arriving, could really mess up the software timing required in B. It gets worse when the exact data received in scenario A may make a huge difference in how much the processor has to do, in the event that the bytes received give some sort of command or require extra processing at the end of variable-length lines of data.

If you look at one scenario or another by itself, the job might even appear possible without interrupts. As soon as you get two or more such scenarios at once, it quickly becomes impossible. Interrupts become imperative. Fortunately, they're pretty easy on the 6502.

1.2 WHAT IS AN INTERRUPT?


Here are some analogies to everyday life, suitable even for the computer-illiterate. The perfect illustration I read in a book (and I don't remember where anymore) used the doorbell for comparison. Suppose you knew one or more guests could be arriving at the door. Polling would be like going to the door often to see if someone was there yet. The analogy with the serial port is to have your program keep asking, "Do you have another data byte yet?" and the answer comes back, "no." Then again "Do you have another data byte yet?" "No." "Do you have..." The better way is for the processor to tell the port, "Ok, let me know when you do," and then quit asking, and go about its business doing something productive.



That's what the doorbell is for. The guests are coming, but you have preparations to make, or maybe something unrelated that you need to do. You go about your business. You only go to the door when the doorbell rings, and the guests are not left waiting for the next time you decide it's time to check the door again. You get more done, and they get quicker response when they ring the doorbell. It works out better for everyone.

The alarm clock is an even more relevant analogy. Without it, you could not sleep soundly and restfully and still make sure you meet your morning obligations. Either you watch the clock for the last few hours of the night (whether that means getting up early or simply lying there making sure you don't fall asleep again), or you risk being late for work or other appointment. (Actually the 65C02 and 65816 do have a WAIT instruction to stop the processor and make it wait for an interrupt. This makes the processor able to respond to the interrupt more quickly because it will not have any instructions in progress that must be finished first. These processors also have a STOP instruction. These might compare to sleeping, so the alarm clock analogy becomes all the more relevant.)

An interrupt is **not** when you decide to take a break from your work. If you're going through the mail and get to a bill and get up to get the check book, that is not analogous to the interrupt, but rather to a branch taken on the basis of what you found in the task you were busy with. There was nothing from the outside that caused you to put down your work and go do something else. Now if you're going through the mail and you hear crying from the back yard because one of your kids got hurt, you put down what you're doing, (hopefully) remembering where you were so you can later pick up again where you left off, and go see what happened. After taking care of the interrupt, however serious or minor it may have been, you return to what you were doing. The only difference it made to the original task was the small delay introduced. Nothing in that original task contained a note telling you to go check the back yard every so often.



Nor is an interrupt when you notice it's time to take lunch. You polled the clock, and compared the time it told you to what time you planned to take a lunch break, and made a decision based on the outcome of the comparison. As the time draws closer, you might "poll" it more and more often. On the other hand, if a bell at work told you it was time for lunch, that would be analogous to the interrupt. You would definitely notice, even if you weren't thinking about it.

An interrupt is not something the program does when it reaches a certain point. (The exception would be the BRK instruction, discussed only briefly later on.) I mention this because some on the forum have talked about doing this kind of thing, not understanding what interrupts are or what they are for. If you want to do something at the end of a loop, fine-- have the program execute it directly or call a subroutine that does it. Interrupts by nature are asynchronous events. You generally don't know where the program pointer will be when an interrupt hits.

Can an interrupt interrupt another interrupt? Sure. Suppose you get out to the yard just about the time little Suzie quits crying, and you see that as usual, it was nothing serious. A Band-Aid would be good just to make her feel like you care; but while you're out there, the phone rings. The phone is seldom important, but it is urgent. Hopefully it's not your wife at the supermarket with a car that won't start. Maybe it's that new customer ready to place a large order that will give your home business a big boost. You don't want him to get the answering machine, so you run in and answer. If it was before the national do-not-call list went into effect, it might have been someone trying to sell you double-pane windows, paint your house, or switch your long-distance service, so you hang up on them (but not before telling them never to call again. Fortunately that's one interrupt we're able to mask through the do-not-call list now since 10/1/03.) Now you get back to Suzie's Band-Aid, finish that up, and then get back to the mail.

1.3 GETTING TO THE NUTS AND BOLTS

The routine that services the interrupt is called the *interrupt-service routine*, or "ISR" for short. Sometimes it is also called an *interrupt handler*. This routine is written like any other, with one exception. Since it can get called at virtually any time, it must remain invisible to the main program. Whatever computer resources it needs to borrow to do its job, it must put back exactly the way it initially found them. It may have an effect on the main program, but it should not cause erroneous operation by pulling a stunt like changing a status flag immediately before the main program was going to use it to decide whether or not to branch, or changing the value in an index register immediately before the main program was going to use it to get a data byte from the desired location. If the ISR borrows a register for use, it must, when it's done, return it to the way it found it. It's not just the courteous thing to do-- it's the only way to keep the computer from crashing as soon as the ISR finishes.

Fortunately, the 6502 automatically takes care of part of this. It puts the return address onto the stack in page 1, and unlike some other processors, also saves the processor status register P on the stack. When a peripheral device needs service and pulls the

interrupt request ($\overline{\text{IRQ}}$) line low, the currently executing machine-language instruction is allowed to finish, and the 7-clock interrupt sequence is started. This sequence is like the execution of an instruction that was not written into the program, but implicitly inserted by the fact that the $\overline{\text{IRQ}}$ line was pulled down to a logic-low level. The interrupt sequence takes two clocks for internal operations, two to push the return address onto the stack, one to push the processor status register, and two more to get the ISR's beginning address from \$FFFE-FFFF (for $\overline{\text{IRQ}}$) or \$FFFA-FFFF (for NMI)-- in that order.

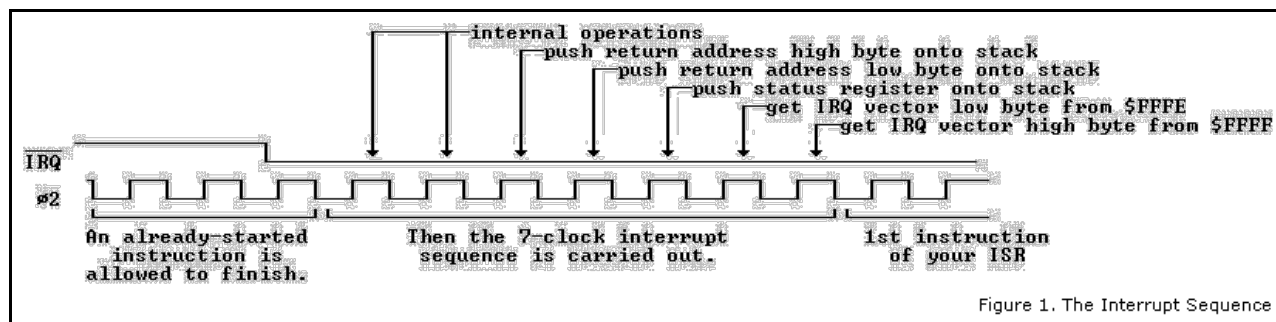


Figure 1. The Interrupt Sequence

Other things are done during the interrupt sequence beyond what is shown in Figure 1. It is important for the programmer to note that the interrupt-disable I flag is set, and that the decimal D flag is cleared on the 65C02 but not affected on the NMOS 6502. If you have an NMOS 6502 and do hex arithmetic in the ISR, you will need to use CLD if there is any chance the program was interrupted at a point where the D flag was set.

The interrupt sequence pushes three bytes onto the stack. First is the high byte of the return address, followed by the low byte, and finally the status byte from the P processor status register.

Remember that on the 6502, the stack grows *down* toward \$100, and the stack pointer register S is decremented immediately *after* each byte is written onto the stack, regardless of which instruction or interrupt condition is producing the stack pushes. For example, if stack pointer register S contained \$EE, the interrupt sequence will put the return address high byte into \$1EE, the low byte into \$1ED, and the status byte into \$1EC, finally leaving S with \$EB. When a byte is pulled off the stack, the stack pointer is incremented *before* the byte is read.

As a parenthetical note, we should mention that the interrupt sequence does not push the return address the same way that JSR does. The JSR does not push exactly the address where the execution should be resumed, but rather the address of the last byte of the three-byte JSR instruction. The RTS instruction at the end of a subroutine must take the address off the stack and increment it before picking up the next instruction. However, the RTI at the end of an ISR does not need to do this incrementing. In other words, PLP RTS is not just a longer way to do RTI. It would actually land you at a different address-- the one *after* RTI would have taken you to.

A second parenthetical note: See also the forum topic, "[A taken branch delays interrupt handling by one instruction.](#)"

The ISR's execution is ended when the RTI instruction (ReTurn from Interrupt, op code \$40), is encountered. After executing the RTI, the processor immediately picks up where it left off in the job it was doing before the interrupt. RTI takes 6 clocks and does the reverse process to put the program counter and the processor status register back. The ISR's RTI is similar to the subroutine's RTS. The primary difference is that RTI restores the status register P too, not just the address to get back to.


RTI's 6-clock process is the reverse of the interrupt sequence. In the first clock, the RTI (op code \$40) is read. The next two clocks are used for internal operations, the fourth for pulling the processor status off the stack, the fifth for pulling the low return address byte off the stack, and the sixth for pulling the high return address byte off the stack. In the very next clock, the op code at the return address is read and the non-interrupt execution resumes.

The interrupt latency on the 6502 is definitely one of the very fastest in the 8-bit world. The interrupt sequence takes 7 clocks and the RTI takes 6, so the overhead is only 13 clocks. Remember that this includes saving and restoring the status register, so you don't have to do that in the ISR. Some comparison of other processors' interrupt performance will be made later. We will also present a little trick afforded by WDC's modern 65C02 which, under special circumstances, allows the interrupt service to be delayed by no more than *one* clock after the $\overline{\text{IRQ}}$ line goes down, instead of the usual 7 to 14 clocks!

If your ISR uses other registers, it will need to save and restore them. If the accumulator is used, you will need to have PHA



(PusH Accumulator) early in the routine, and PLA (PulL Accumulator) somewhere right before the RTI. The CMOS 6502 (65C02) can directly push and pull the X and Y index registers with instructions PHX, PLX, PHY, and PLY. If you're using the NMOS 6502, you won't have those instructions available and will have to do a work-around. Pushing X on the stack requires using A, so you have to push A first to avoid permanently losing its value (whether your ISR needs A or not), then transfer X to A with TXA, and do another PHA. Pulling it off the stack is the reverse, using PLA and TAX, before restoring the original A with PLA again right before exiting the ISR.



If the ISR needs a variable that the main program might also need, you will have to save that as well. LDA, PHA will do the job. Just remember that whatever you push onto the stack, you must pull off in reverse order, since the stack by definition is a last-on-first-off memory, the opposite of FIFO (First-In-First-Out memory).

So suppose your ISR needs to use A and X. It will need:

```
ISR:   PHA
       TXA
       PHA
       (actual programming to service the interrupt goes here)
       PLA
       TAX
       PLA
       RTI
;-----
```

If you have a CMOS 6502, you can shorten that to:

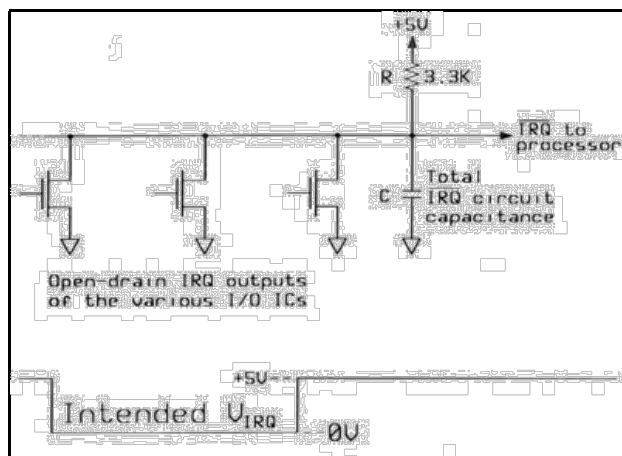
```
ISR:   PHA
       PHX
       (actual programming to service the interrupt goes here)
       PLX
       PLA
       RTI
;-----
```

Example: If your ISR only needs to increment a 16-bit counter in memory you have called CNT, it won't need to use any processor registers. However you do still need to turn off the individual interrupt. For the sake of discussion, let's say the interrupt came from T1 (timer 1) timing out in a 6522 VIA. Reading the timer's low counter byte clears the interrupt so the processor doesn't immediately head back into the ISR again as soon as it finishes. The entire ISR would be:

```
CNT_ISR: BIT   VIA_T1CL ; Turn off interrupt early. (More on that below.)
          INC   CNT      ; Increment the low byte of the variable.
          BNE   isr1$    ; Branch to end if the low byte didn't roll over to 00.
          INC   CNT+1    ; Otherwise increment high byte of variable.
isr1$:   RTI            ; Exit the ISR, restoring the previous processor status.
;-----
```

The INC instruction increments the memory location without changing A, X, or Y. The BIT instruction is used here to read the timer 1 counter low byte without putting it into a register. The processor status register P gets modified, but remember the IRQ sequence and RTI instruction take care of saving and restoring the status so we don't foul up the main program's operation immediately after the ISR runs. Bracketing your ISR with PHP and PLP would be redundant and a waste of execution time. (Again, remember also that the interrupt-disable bit I gets set in the interrupt sequence after P is pushed, so doing an SEI would also be redundant, and a waste of execution time.)

Now I must mention a hardware consideration here. The ISR above first tells the VIA that it can quit asserting the interrupt. There's a reason for doing this first. Most of the 65-family I/O ICs have open-drain $\overline{\text{IRQ}}$ outputs, intended to be connected to each other and to the processor's $\overline{\text{IRQ}}$ input. Any of the I/O ICs can pull the line low, but the only way it will go up is with a pull-up resistor pulling the line up to the positive power supply voltage. The recommended resistor value is usually 3K. Well, since the $\overline{\text{IRQ}}$ line has some amount of capacitance to ground and other nets, that capacitance multiplied by the pull-up resistor value gives you a time constant. If you put the $\overline{\text{IRQ}}$ -clearing instruction(s) at the end of the ISR, it's possible that there won't be enough time for the resistor to charge that capacitance up to a valid logic-high state before the RTI is done executing. The result will be that as soon as the ISR is done and the old status



register flags are restored with the interrupt-disable bit clear once again, the processor is again open to handle another interrupt, so it will respond to the still-low \overline{IRQ} line because the \overline{IRQ} has not had enough time to float back up. Remember the rise time with the resistor is much slower than totem-pole-type logic outputs. If you only have one interrupt source on the \overline{IRQ} line, a 3K pull-up resistor, and only 1MHz clock speed, you can normally get away with clearing the interrupt near the end so you don't need to worry about it; but turn up the clock speed and put more \overline{IRQ} sources on the line, etc., and you can get into trouble. I mention it because I've been bit by it.




Figure 2. The Wire-OR concept, showing rise and fall times.

While most 65-family I/O ICs have open-drain \overline{IRQ} outputs, WDC's 65C22 VIA (Versatile Interface Adapter) does not. That means that to feed more than one WDC 65C22's \overline{IRQ} output into the processor's \overline{IRQ} input, you'll need an AND gate arrangement, so that if any one of the \overline{IRQ} 's goes low, the processor's \overline{IRQ} input is taken low. The AND gate may make for more parts, but the faster resulting rise time takes care of the potential problem mentioned above, especially with their faster processor speeds of 10MHz and higher.




Figure 3. AND Gate arrangement, with WDC VIAs and other open-drain \overline{IRQ} outputs.

You may have already thought of the question of how to know which device caused the interrupt if several are connected to the \overline{IRQ} line at once. Well, no chip capable of generating an interrupt will do so without your program having set it up to do it; so you already know that an interrupt will not have come from a chip that your program didn't set up to generate interrupts. If that still leaves two or more chips that could have caused an interrupt, we do what is called "polling". This is not the same kind of polling we talked about earlier, analogous to checking the door frequently to see if someone was there. This one is more like having a doorbell that sounds the same whether the button was pushed at the front door or at the back door. You might choose to first check the door that's most likely to have someone ringing the doorbell. Alternately, you might choose to start by checking the one most urgent to answer. For example, if you're expecting an important guest but there's a possibility that your seven-year-old is pushing the button at the back door because he locked himself out of the house again, you might still decide to check the front door first to keep the important guest from waiting when the seven-year-old should have remembered to unlock the door before he went out and shut it behind him.

Most 65-family I/O ICs use the highest bit (MSB) of their status registers as a master interrupt bit. Simply reading the status register transfers its high bit to the N flag in the processor's status register, and you can branch on that with BMI or BPL. LDA VIA1_STATUS will put bit 7 of the VIA's status register into the processor's N flag. One of the functions of the BIT instruction however is also to do the same thing, but you don't have to touch the accumulator's contents to do it. Take the example:

```
ISR: BIT   VIA1_STATUS    ; Check 6522 VIA1's status register without loading.
      BMI  SERVICE_VIA1  ; If it caused the interrupt, branch to service it.
      BIT  VIA2_STATUS    ; Otherwise, check VIA2's status register.
      BMI  SERVICE_VIA2  ; If that one did the interrupt, branch to service it.
      JMP  SERVICE_ACIA   ; If both VIAs say "not me," it had to be the 6551 ACIA.
;-----
; (Don't forget that the last ISR instruction to be
;      executed must be RTI, not RTS.)
```

This ISR assumes there are three possible sources of interrupts. (The two 6522's and a 6551 are what [Daryl Rictor's SBC-2 board](#) has.) The ISR first checks with the #1 6522 VIA. If that one is not the one that caused the interrupt, check the second VIA. If that one didn't either, the assumption is that the ACIA did it, since there's nothing else left.

Now if two devices caused interrupts at the same time, you'll poll the higher priority one first. Finding that it did indeed need service, you will service that one before the lower-priority one.

If the routine to service the first chip ends in RTI, the second chip will not have gotten serviced yet. That's ok though; because as soon as the RTI is finished, the second interrupt source will still be holding the \overline{IRQ} line low, so the ISR will immediately be entered again. (Remember that the 6502's \overline{IRQ} input is level-sensitive, not edge-sensitive.) The first chip will normally not have generated another interrupt again so soon, and the first poll will test false so the program counter will drop through to test and service the second chip.

The RTI followed immediately by another \overline{IRQ} sequence above may seem like a waste of time. What you could do is have SERVICE_VIA1 end by jumping back to the third line of ISR above, so VIA2 will get polled before the RTI, and have SERVICE_VIA2 jump back to the fifth line above, so the ACIA would get polled too. However, all the extra polling is a waste of time too. It would normally be best to end SERVICE_VIA1 and SERVICE_VIA2 with RTI just like SERVICE_ACIA, instead of coming back to check for more possible interrupt sources before exiting the ISR.

An exception would be where you expect two or more interrupts to regularly hit at once. In that case, you might want to

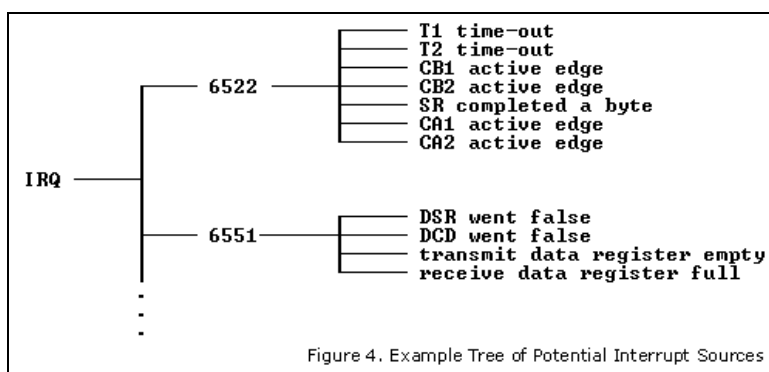
re-enable interrupts (using CLI) in the ISR as soon as the most urgent part of one interrupt's servicing is finished (including turning off the assertion of that interrupt). This is so that another interrupt can have its most urgent part serviced sooner too, instead of having to wait for the completion of the first interrupt's service.

One situation that is usually undesirable is where an interrupt source interrupts its own service; ie, you don't have enough processing power to service the interrupts as fast as a source is generating them. This situation is precarious at best. This might be remedied in various ways, from making your software more efficient, to changing how your software sets up those ICs so you slow them down. This may include reducing the baud rate on a serial link if there's no hardware handshake, lengthening a timer's time-out period, or reducing the rate that something external produces interrupts. The only other solution is just to increase the processor speed.

If you know that a regular, predictable interrupt from a source such as a timer will be hitting often enough to be pushing the limits of how fast you can service it, you better count cycles in the ISR (not forgetting to include the IRQ & RTI overhead and the time to finish an instruction that was already executing when the interrupt hits), and make sure that the cycle count, divided by your $\phi 2$ rate, leaves you enough time between interrupts to service them completely.

Remember that the routines branched to from the piece of code above must save the accumulator on the stack and later restore it (assuming they will in fact use the accumulator), as well as X, Y, and possibly certain variables that the main program might also be using. If all three of the routines branched to above needed to push the accumulator onto the stack, it would be more economical in terms of program memory to just do the PHA immediately at label ISR instead of in three different places. The same goes for other things they may all need to save. If the IRQ line was pulled down, at least one of these will need to run, so saving the registers up front will not be a waste of time.

After determining which IC caused the interrupt, you might need to see which of two or more possible causes within that IC set the interrupt bit. Testing individual bits of the I/O chip's status register will reveal that. Enabling the chip to interrupt typically means enabling particular sources of interrupts within the chip. Obviously there's no point in testing for an interrupt that was not enabled. For example, you may have enabled a 6522 VIA to interrupt when timer 1 (T1) times out, but not when timer 2 (T2) times out, nor when the serial port's shift register has completed shifting a byte in or out, nor any of the other interrupt possibilities. In that case, you don't need to see which bit in the interrupt flag register (IFR) is set, because there is only one possibility. So if the VIA's status register's bit 7 is set, you know T1 timed out.



If there are two or more interrupt sources enabled in the chip, then you'll have to check bits, again checking for the higher priority one first in case there were two sources of interrupt in the chip at the same time. You can check with the BIT instruction after putting the bit mask in A, or read the chip's status register into A and use the AND instruction. Either case will be followed by a conditional branch instruction, probably BNE or BEQ. Just a note about the BIT instruction-- It does not matter what's in the accumulator if you only want to test bit 7 or bit 6 of the memory location (in this case an I/O chip's status register). You *will* need to load a bit mask into A to use BIT to test other bits of that memory location however.

An applications book pointed out that you could, in the case of the 6522 for example, load the IFR (interrupt flag register) into the accumulator, and shift left repeatedly, each time branching on the N flag results, like this:

```

LDA VIA_IFR          ; Get the set of interrupt flags, and
BPL <test next device> ; skip down to test next device if this 6522 is not
ASL                  ; the one that requested service. Else shift the T1
BMI SERVICE_TIMER_1  ; IFR bit into the bit-7 position to test...
ASL
BMI SERVICE_TIMER_2
(etc)
  
```

Note however that particular interrupt flags in the IFR might be set without generating an interrupt. For example, bit 2 may be set meaning that the serial port has finished shifting a byte in or out; but if you intentionally are not using interrupts for your serial port operations, you would not have had the corresponding bit in the IER (interrupt enable register) set, meaning that the

"shift completed" status was not what pulled the $\overline{\text{IRQ}}$ line down. To avoid unintentional branching, AND the IFR with the IER:

```
LDA VIA_IFR          ; Get the set of interrupt flags, and
BPL <test next device>
AND VIA_IER          ; zero those that were not allowed to pull  $\overline{\text{IRQ}}$  down.
ASL
BMI SERVICE_TIMER_1
ASL
BMI SERVICE_TIMER_2
(etc)
```

So far the only I/O ICs mentioned here are the 6522 VIA and the 6551 ACIA. There have been more than 20 other I/O ICs in the 65xx family alone, and certainly hundreds of other ones that are not in the 65 family but nevertheless can be interfaced to the 6502.

1.4 SETTING UP AN INTERRUPT

So how do you set up an interrupt on the 6502?

The 6502 and all its variants power up with the interrupt disable bit (I) set, meaning the processor won't initially respond to anything that pulls the $\overline{\text{IRQ}}$ line low. It will ignore the $\overline{\text{IRQ}}$ line. For interrupt service, you must:

1. make sure the proper program (ISR) is in memory, ready to service the interrupts. The IRQ (or NMI) vector must point to the beginning address of this routine.
2. set up the IC or ICs you want to enable to produce the interrupts
3. clear the interrupt-disable bit in the processor status register, with CLI, so the processor will respond to interrupts.

Let's go back to the time counter used earlier, and assume for the sake of simplicity that it's the only interrupt on the μP 's $\overline{\text{IRQ}}$ input. Step 1 requires the ISR that was already given above:

```
CNT_ISR: BIT    VIA_T1CL ; Turn off interrupt early (as discussed above).
          INC     CNT     ; Increment low byte of variable.
          BNE    isr1$   ; Branch to end if the low byte didn't roll over to 00.
          INC     CNT+1   ; Otherwise increment high byte of variable also.
isr1$:  RTI          ; Exit the ISR, restoring the previous processor status.
;-----
```

The *interrupt vector* should point to the address at label ISR. So what is an interrupt vector?? It's a designated address in memory whose contents tell the processor the ISR's beginning address. The 6502 will read addresses \$FFFE and \$FFFF to get the address of the ISR. (Actually, my books are not consistent as to whether \$FFFE-FFFF *are* the vector, or *contain* the vector. It doesn't really matter much though-- the important part is understanding what the processor does.) A section at or near the end of your source code might look like:

```
.ORG    $FFFA      ; Make the assembler start this section at addr $FFFA.
NMIVector: .WORD   NMI_ISR ; Make the NMI vector point to NMI ISR if you have one.
RESVEC: .WORD   RESET    ; Make the reset vector point to the reset routine.
IRQVEC: .WORD   CNT_ISR  ; Make the IRQ vector point to the timer counter ISR.
```

Here we're making this counter to be the IRQ interrupt service, and assuming the VIA's $\overline{\text{IRQ}}$ output will be connected to the μP 's $\overline{\text{IRQ}}$ input and not the NMI input. More on that later. Suppose CNT_ISR (the routine above) started at address \$E06A. Address \$FFFE of the IRQ vector would contain \$6A (the low byte of \$E06A), and address \$FFFF would contain \$E0 (the high byte of \$E06A). Since the reset vector at \$FFFC-FFFF must in most cases be in ROM so they're there when the computer first powers up, you will normally have ROM in that part of the memory map, so the interrupt vectors will also be in ROM. If you want to be able to change it easily as different programs are loaded, you might want this vector in ROM to point not directly to the ISR, but to a place in RAM that says JMP xxxx, where the xxxx is the address of the ISR you want executed. This JMP xxxx instruction must be stored there by the program before interrupts are enabled. It adds three more clocks to the interrupt latency, but allows you to effectively change the vector without changing the ROM. Another way to handle it would be to just make sure the ISR itself always starts at that same address.

For step 2, let's set up timer 1 (T1) in a 6522 VIA to time out every 50,000 ϕ_2 clocks and produce an interrupt every hundredth of a second (at 5MHz ϕ_2 rate):

```
TMR_SETUP: STZ    CNT          ; Initialize the count that will be incremented by
          STZ     CNT+1        ; the ISR at every time-out of T1.

          LDA     #$4E         ; Put $C34E (50,000-2) in the VIA's timer 1 counter
          STA     VIA_T1CL     ; low and high bytes. Note: you must write to the
          LDA     #$C3         ; counters to get T1 going. After that, you can
          STA     VIA_T1CH     ; write to the latches. $C34E will make T1 time out
                                ; 100 times per second at 5MHz.
          LDA     VIA_ACR      ; Clear the ACR's bit that
          AND     #01111111B   ; tells T1 to toggle PB7 upon time-out, and
```



```

ORA  #0100000B ; set the bit that tells T1 to automatically
STA  VIA_ACR    ; produce an interrupt at every time-out and
                ; just reload from the latches and keep going.

LDA  #1100000B
STA  VIA_IER    ; Enable the T1 interrupt in the VIA.

```

T1CL above stands for "timer 1 counter low byte", T1CH is "timer 1 counter high byte," ACR is "auxiliary control register", and IER is "interrupt enable register." The 6522 data sheets will tell the significance of individual bits in the various registers and how to calculate the number to put in the T1 latches to make the timer run at the desired rate.

Step 3 above is pretty self-explanatory. You might want to put the CLI at the end of the TMR_SETUP routine above.

What's above is a simpler version of the real-time clock (RTC) that I have on my workbench computer for time of day (with .01 second precision) and calendar. It slows the whole system down by about 1%, including looking to see if an alarm is due. An alarm would tell the computer it's time to do a particular job. I have VIA1 connected to NMI instead of IRQ though, for a couple of reasons. It's the only thing I have on NMI, so it requires no polling overhead of its own, and neither does it add any polling overhead to the IRQ ISR. Being on NMI, it cannot be affected by SEI (the set-interrupt-disable instruction) or a super-long-running IRQ ISR. (The latter may be a moot point since ISRs virtually never take anywhere near the 50,000 clock cycles we get between T1 interrupts as set up in the above fashion.

The RTC can of course be stopped if it would cause too much jitter in another job being done in the foreground. Subroutines to turn the timer counting off and on might look like:

```

TRN_CNT_OFF: LDA  #0100000B
              STA  VIA_IER
              RTS
;-----
TRN_CNT_ON:  LDA  #1100000B
              STA  VIA_IER
              RTS
;-----

```

Up until the last couple of paragraphs, we have only talked about interrupts on the 6502's **IRQ** input. There is also an **NMI** (non-maskable interrupt) input on the 6502. Note that 6502-based microcontrollers may have many interrupt inputs from various onboard I/O blocks, so each interrupt source can have its own vector and polling is greatly reduced. When it's not in a microcontroller chip however, the 6502 microprocessor by itself only has the two interrupt inputs, **IRQ** and **NMI**, and their two corresponding vectors. More about NMI later.

Interrupts on the **IRQ** input can be masked; ie, you can have the processor ignore them. Sometimes a program may need to do something that you don't want interrupted. It may, for example, be that you're doing some operation on a section of memory and the ISR could foul things up if it were allowed to alter anything in that section of memory before you're done. Another possibility is that the timing on an I/O operation may need to be tightly controlled for an instant, and you can't afford to have an ISR jumping in during that time. To mask IRQ interrupts, use the SEI (SEt Interrupt-disable bit) instruction. To make the processor able to respond to IRQ interrupts again, use the CLI (CLear Interrupt-disable bit). Each of these takes only two clocks to do its job of setting or clearing the I bit (bit 2) in the processor status register P.

Note that neither CLI nor SEI clears or acknowledges any interrupts! It does not affect anything outside the processor itself. Also note that the 6502 does not have any interrupt-acknowledge pin. Interrupt conditions are cleared in the interrupting IC by reading or writing particular registers in it per its data sheet. This is also the only acknowledgement.



The 6502's **IRQ** input is level-sensitive. Anytime this input pin is low and the interrupt-disable bit is clear, the processor will finish up the currently executing instruction and begin the interrupt sequence. Again, this sequence consists of pushing the return address and the processor status onto the stack, and getting the interrupt vector from address \$FFFE-FFFF and beginning execution at the address pointed to there.

As you can imagine, there would be trouble if the interrupt sequence were performed over and over as long as the **IRQ** input were held low. The stack capacity of 256 bytes (all of page 1, from address \$100 to \$1FF) would be exceeded after pushing the return address and the processor status a maximum of 85 times; but no more useful code could ever get run anyway so the stack overflow would not matter.


What saves the day is that after the 6502 pushes the processor status register P onto the stack, it sets the I (interrupt-disable) bit

so further interrupt requests will be ignored until the status is restored at the RTI, or until you use the CLI instruction to explicitly clear the I bit inside the ISR. Note: The ISR should not clear the I bit until after it has found and turned off the source of the interrupt that caused the ISR to be entered. Otherwise we will again be back to the situation of the computer crash caused by endless interrupt sequences as just mentioned.

"Turning off the source of the interrupt" does not mean disabling all interrupts, or even all possible sources of interrupts from the particular IC. In the counter ISR example above, the BIT VIA_T1CL instruction reads the VIA's timer 1 counter's low byte, which makes the VIA quit asserting the interrupt until the next time the timer times out-- ie, when its count reaches 0 and gets reloaded from the value you stored in the latches (\$C343 in the example above). Notice that we didn't write to the VIA's interrupt-enable register again (IER). The interrupt is still enabled, but it ceases to be asserted and will not be asserted again until the next time the count gets decremented through 0. The data sheets tell what action makes any given interrupt release the IRQ line so it can float up again.

Most interrupt sources will be connected through the processor's IRQ input. It is level-sensitive for a good reason. Suppose an interrupt occurs, and the processor carries out the interrupt sequence and enters the ISR. Now before that particular interrupt is turned off, suppose another interrupt occurs, also holding the IRQ line low. When the servicing of the first interrupt is finished, the IRQ line continues to be held low, so there will be no further edges. Having the IRQ input to be edge-sensitive would not be appropriate. The continued low level on the IRQ line is the only way for the processor to know that there is another interrupt needing service. If the ISR has no CLI instruction and does not check for further interrupts before finishing (as is normal, as discussed earlier), then the interrupt sequence will be carried out again right after the RTI instruction.

There is a possible situation that should be mentioned, regarding "ghost" interrupts. If a chip produces an interrupt during an instruction to disable that same interrupt, the chip will pull the IRQ line down quickly and release it just as quickly. If it takes the pull-up resistor a couple of clock cycles' time to get the line up to a logic 1 again, that's enough time for the processor to respond to the interrupt and begin the interrupt sequence. Remember that the chip may also need a cycle or two to process the interrupt-disabling command byte you just wrote to it. But by the time the ISR checks to see if this chip caused the interrupt, the chip's interrupt flag register will say "none here"-- like ringing the doorbell and running. I had this situation once where I had a VIA's T1 producing 50,000 interrupts per second, and the program was turning the interrupts off and on to produce one-tenth-second bursts of interrupts. If I remember correctly, each time one of the interrupts hit, the ISR would send the next sample of a waveform to the D/A converter to produce an analog signal. This signal went in one-tenth-second bursts, with a tenth of a second between bursts. Sometimes, turning off the burst produced a ghost interrupt.



There are different ways this could be handled. The best one is probably to use the SEI instruction before disabling the particular interrupt, and the CLI again after the potential for generating the interrupt is gone.

```
SEI
JSR  TRN_CNT_OFF
CLI
```

(The SEI and CLI could even be made part of subroutine TRN_CNT_OFF.)

Another way would be to have the ISR itself disable the particular interrupt source, since the processor won't normally respond to interrupts during the ISR anyway. This could make the ISR too long though, if it has to check every time and see if it is time to disable the particular interrupt. A quick and dirty way to handle it might be that if you know this could happen and the ISR does not find any cause for the interrupt, you could just have it exit without doing anything, with the expectation that this course of action will be adequate.

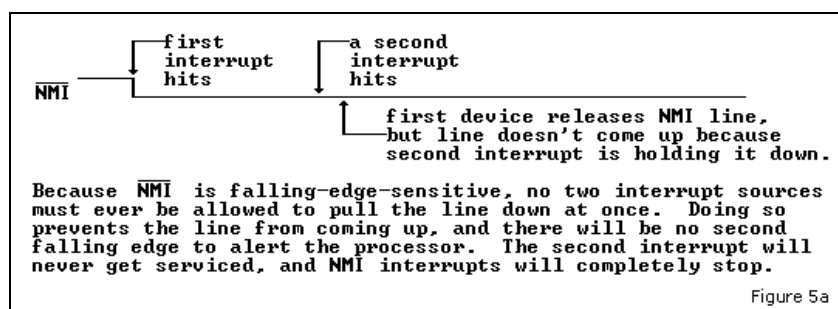
An ISR's job should normally be kept as small as practical in order to maximize overall interrupt performance and make sure all deadlines are met, particularly if there is more than one interrupt source. Any work that is not time-critical should be left for the main program to do. This may simplify your programming job too.

Someone was telling me about a project where a machine's position counter had to be read about every millisecond, regardless of anything else happening at the moment. In other words, there was a deadline that had to be met, with no exceptions, about every millisecond. Any less often could mean that the counter could roll over and information would be lost, and the product being machined could be ruined. The position information was to be displayed, and if I understood correctly, may also be used

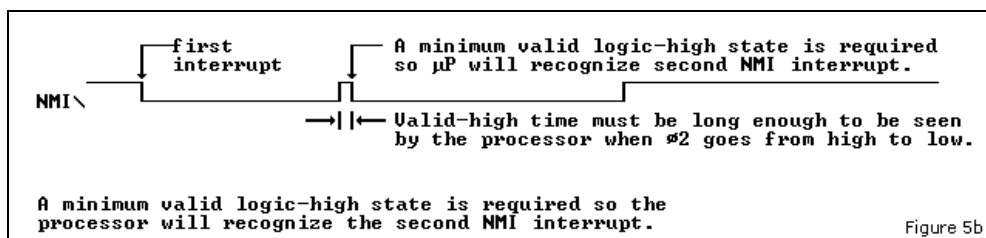
to control the motors to achieve a desired position. Here the ISR probably only needs to read the position counters. The motors' commands would not have to be changed as often, and there's certainly no point in updating the display more than about ten times per second. The jobs of sending commands to the motor controllers and updating the display should be left for the main program, simplifying the ISR and making sure the part that absolutely *must* be done every millisecond *is*. Then it wouldn't matter if any of the other parts take over a millisecond to carry out, because they won't be interfering with the part that cannot be compromised.

2.1 NMI: THE NON-MASKABLE INTERRUPT

Now back to the 6502's other interrupt input we've hardly mentioned so far, the **NMI**. This one is not level-sensitive, but edge-sensitive. Since it is not maskable, making it level-sensitive would have brought about the same endless-interrupt-sequence disaster mentioned earlier. However, the fact that it is edge-sensitive limits it to generally only one interrupt source. If there's more than one, steps will have to be taken to make sure that no two sources can pull it down at once, or at least that there is no way that one source can get unnoticed by the ISR and be left unserved, thereby holding the NMI line low such that no more edges can occur. See Figure 5a.



Furthermore, you need to make sure that some minimum amount of time between low states be observed so that the processor has time to see that the **NMI** voltage actually went up and back down. The **NMI** line must be at a valid logic high state long enough to be seen by the processor at the time of one ϕ_2 falling edge. It can be low again for the next ϕ_2 falling edge, and the processor will recognize the new interrupt. These ϕ_2 falling edges do not have to be at any particular part of an instruction in progress. See Figure 5b.



Since it is not maskable, an interrupt on **NMI** can cut in on the servicing of an **IRQ** interrupt, even though the interrupt-disable I flag is set. Note also that if an **NMI** and an **IRQ** hit at the same time, the **NMI** has the higher priority and will get serviced first.

We're often told to reserve the **NMI** for something drastic like power going down; but in most systems the people on 6502.org are making, what happens in the last milliseconds before power is gone is of no concern. If you have a system that remembers things when it's off, it probably has batteries and can turn itself off in an orderly fashion. Otherwise, if you accidentally pull the power cord, there's no time to store anything useful on a disc anyway.

Scott Schidester on the 6502.org forum said on Jan 6, 2003, "On my little SBC project I have a 'panic' button attached to **NMI** whose handler basically drops into a monitor prompt..." Good idea. Just don't forget to use an RC and schmitt-trigger gate for debouncing so you don't overrun the stack with **NMI** return addresses and status bytes, which could happen if you don't deliver a single, clean **NMI** edge. You can also use such a panic button or "ABORT" button to do a program reset to regain control of a runaway program without delivering a reset pulse to the I/O ICs. There have been times I've wanted to do this while, for example, developing an application in order to keep something interrupt-driven running more-or-less uninterrupted by the semi-reset.

If the 6522 VIA used for the real-time clock (RTC) goes on the 6502's **NMI**, polling for interrupt sources on **IRQ** is simplified, and the RTC never misses a beat when other interrupts hit or are masked.

The following example of using a 6522 VIA's timer 1 for the RTC is similar to what is in my workbench computer code. At 5MHz, the longest round interval T1 can give is ten milliseconds, so I set it up to produce an interrupt precisely that often--10ms. Even if you don't need to know the time of day or the date, the RTC is useful for other things like timing how long you

The whole NMI ISR below (called INCR10ms), if you use the whole thing, is rather long; but note that only 10 instructions total get executed in 98.6% of the NMI occurrences. For midnight New Year's where the carry ripples all the way into the incrementing of the year, 45 instructions get executed. This makes for an unusually long ISR. If this rare occurrence forces an unacceptably long delay on other ISRs (those coming from the IRQ input), you might want to re-enable the IRQ interrupts (using CLI) after the incrementing of the seconds. Any routine (including ISRs) that need the time and date variables should read the needed ones until it gets the same result twice in a row anyway, in order to avoid invalid results from the set of variables being updated between the times the routine begins and finishes reading them, or in this case, their being only partially updated. Interestingly, allowing an IRQ interrupt to interrupt the NMI ISR would mean you would have ISRs nested at least two levels deep; but the roll-over of the year (or month, or hour) itself will not corrupt the end result. The next 99 NMI occurrences will only run ten instructions each anyway, and the rest of the time and date updates don't need to be finished before the next second is up.

25/09/2016, 13:03

```

        INC    SECONDS          ; increment the seconds.
        LDA    SECONDS
        CMP    #60              ; See if seconds carries to another minute.
        BMI    end_NMI          ; If not there yet, skip the rest of this.
        STZ    SECONDS          ; Otherwise zero it,
                                ; and go on to
        INC    MINUTES          ; increment the minutes.
        LDA    MINUTES
        CMP    #60              ; See if minutes carries to another hour.
        BMI    end_NMI          ; If not there yet, skip the rest of this.
        STZ    MINUTES          ; Otherwise zero it,
                                ; and go on to
        INC    HRS              ; increment the hours.
        LDA    HRS
        CMP    #24              ; See if hours carries to another day.
        BMI    end_NMI          ; If not there yet, skip the rest of this.
        STZ    HRS              ; Otherwise zero it,
                                ; and go on to
        INC    DAY              ; increment the day.

        LDA    MO               ; Now the irregular part.
        CMP    #2               ; Is it supposedly in February?
        BNE    notfeb           ; Branch if not.

        LDA    YR               ; For Feb, we have to see what year it is.
        AND    #11111100B       ; See if it's leap year by seeing
        CMP    YR               ; if it's a multiple of 4.
        BNE    notfeb           ; Branch if it's not; ie, it's a 28-day Feb.

        LDA    DAY              ; Leap year Feb should only go to 29 days.
        CMP    #30              ; Did Feb just get to 30?
        BEQ    new_mo           ; If so, go increment month and re-init day to 1.
        PLA                    ; Otherwise restore the accumulator
        RTI                    ; and return to the regular program.

notfeb: PHX                    ; Save X for this indexing operation.
        LDX    MO               ; Get the month as an index into
        LDA    MO_DAYS_TBL-1,X ; the table of days for each month to increment,
        PLX                    ; and then restore X.
        CMP    DAY              ; See if we've reached that number of days
        BNE    end_NMI          ; If not, skip the rest of this.

new_mo: LDA    #1               ; Otherwise, it's a new month. Put "1" in
        STA    DAY              ; the day of month again,
        INC    MO               ; and increment month.
        LDA    MO
        CMP    #13              ; See if it went to the 13th moth.
        BNE    end_NMI          ; If not, go to end.
        LDA    #1               ; Otherwise, reset the month to 1 (Jan),
        STA    MO

        INC    YR               ; and increment the year.
end_NMI: PLA
        RTI
;-----

```

If another part of your code runs often enough to take care of everything past the hours, you can make the ISR shorter. For example, if a display routine will get run at least once per hour all the time the computer is on, you might be able to have it call a routine that takes care of everything past the incrementing of the minutes.

The NMI ISR above was done in two sections, `cs_32` and `year-month-day-hour-minute-seconds-hundredths`, because it was the easiest way to have both the time-and-date-relevant records and the absolute "ticker" of hundredths of seconds. Incrementing both in the NMI ISR was easier and gave less overhead than having to convert each time you need the one you don't have.

Mine actually has even more to it, comparing `cs_32` to the 32-bit timer value for the next alarm to come due. (5/17/14: There's a description and code in the [Alarms](#) section of my [article](#) on simple methods for multitasking without a multitasking OS for systems that lack the resources to implement a multitasking OS, or where hard realtime requirements would rule one out anyway.) If the times match, it sets the high-level interrupt status for my Forth system to service the alarm. This service can consist of anything the computer can do, from acting like an alarm clock to performing a complex operation of gathering data and controlling equipment. The alarm service routine can in turn also take care of setting up the next alarm. Adding the alarm part only slows things down by .034% if an alarm is pending, and .014% otherwise.

Since the RTC ISR can cause some jitter in other timed events, I sometimes turn it off by using `RTC_OFF` (shown above). "[Jitter](#)" is when the timing of a regular event gets skewed. When I've done sampling for audio digital recording or playback at anywhere up to about 50,000 samples per second using interrupts from another VIA's T1 on the processor's IRQ input, I turn off this RTC. Otherwise some of the audio samples get taken care of a little late instead of precisely at the right times. Another way

to do it would be to move the audio sampling clock to the NMI, and the RTC to IRQ.

When you need to read more than one byte of the time variables, make sure you read them twice in a row and get the same result. If it's not the same, read it again. If not on the first try, the second try should usually give you the same result twice. This is important because if the interrupt hits and increments the time while you're reading it and are between bytes, you may get an invalid result. Getting the same result twice guarantees it's valid.

2.2 BRK: THE SOFTWARE INTERRUPT

There is also a software interrupt, the BRK (break) instruction, which goes through the motions of a hardware interrupt, without anything having pulled the $\overline{\text{IRQ}}$ or $\overline{\text{NMI}}$ line down. The vector used is at \$FFFE-\$FFFF, the same one used by IRQ. BRK was mainly used for patching code in PROMs back when re-assembling and programming was a long, slow process, and each iteration could be quite expensive if the PROM was not erasable. The BRK instruction seems to mostly have outlived its usefulness. There is some use for it in multitasking operating systems, but the 6502 is not generally very well suited for that anyway. The 65816 as other capabilities that make it better suited to multitasking, relocatable code, etc..

If you do use BRK, the ISR will normally need to test the B bit (the "break" bit, bit 4) in the *stacked record* of the processor status register to determine if the interrupt was caused by the BRK instruction. The ISR cannot test it in the status register itself by doing PHP, PLA, AND #00010000B, BEQ/BNE, because doing this will always make it appear to be set. From inside the ISR, you must do PLA PHA (to take the status byte record that was stacked at the time of the interrupt and to copy into the accumulator), followed by AND #00010000B and BEQ/BNE. So again-- the test must be preceded by PLA PHA, not PHP PLA. The latter will not work, since it will always show bit 4 set.

BRK does set the interrupt-disable I flag like an IRQ does, and if you have the CMOS 6502 (65C02), it will also clear the decimal D flag.

Note that BRK, although it is a one-byte instruction, needs an extra byte of padding after it. This is because the return address it puts on the stack will cause the RTI to put the program counter back not to the very next byte after the BRK, but to the *second* byte after it. This padding byte can be used for a signature byte to tell the BReaK interrupt routine which BRK caused the particular software interrupt.

We should mention here that one of the NMOS 6502 bugs is that if an NMI hits during a BRK instruction, the BRK interrupt will not get executed. This and all NMOS 6502 bugs have been fixed in the CMOS 65C02-- giving more reasons to switch to the CMOS version.

2.3 WAI: FASTER INTERRUPT SERVICE ON THE 65C02

WDC's 65C02 has a "wait" instruction, WAI. This allows a special case of ultra-fast IRQ interrupt service.

In the earlier doorbell comparison, the guests, after ringing the doorbell, had to wait for you to put down what you were doing and get to the door, which introduced a small delay before they would see the door open. The comparison now is that you have already put your work down and gone to the door so you're right there ready to open it immediately upon hearing the bell.

If the main part of the work can be paused until the next interrupt, you can set the interrupt-disable bit I-- yes, *set* it-- execute WAI, and have the very next instruction to be the beginning of the ISR. There will be no jumping through vectors; and since you know exactly where the program pointer will be when the interrupt hits, you will not necessarily need to save any registers your ISR uses. If you do, you can do it *before* the interrupt.

The WAI guarantees that the processor will not be in the middle of executing another instruction when the IRQ line is pulled down, so we can eliminate that part of the latency. The other part of the latency, the 7-clock interrupt sequence, gets eliminated by the fact that we have used SEI to disable the normal IRQ operation, so the IRQ will only have the effect of re-starting the processor and making it continue on with the next instruction instead of taking the vector. And since we don't take the vector, we won't use RTI at the end either.

Here's the idea. The LDA and STA instructions were selected only arbitrarily for the example. The xxx just represent the continuation of code execution after the interrupt service is finished.

```

STA VIA1IER      ; Pre-interrupt code finishes here.
SEI              ; Disable interrupts if not already done.
WAI              ; Put processor into pause mode.
LDA VIA1PB       ; First instruction of ISR goes here.
.               ; (Notice the code is straight-lined.)
.               ; Service the interrupt.
.
xxx              ; End of ISR moves right into the next thing
xxx              ; for the computer to do, without using RTI.
```

3.1 INTERRUPT SUPPORT FOR RS-232 RECEIVE

As promised in the beginning, here's the program example of receiving data by RS-232 in the background through a 6551 ACIA, driven by interrupts, and putting the data into a buffer that holds it until the program running in the foreground is ready for it. (There is an RS-232 primer at <http://wilsonminesco.com/RS-232/RS-232primer.html>.)

The reason for having the ACIA generate an interrupt each time a byte is finished coming in is that this allows the computer to do other useful things in the foreground while it's waiting for the ACIA to deliver the next byte. We will put the data into a 256-byte buffer, for a few reasons: 1. If the foreground program is not ready for more data yet, we can make better use of the serial port time by keeping the data coming anyway, and put the data into this buffer. 2. When the program is ready for more data, it may take a chunk of it faster than the serial port would be able to deliver it. If there's already some data in the buffer, the program won't have to wait for it. 3. If your 6502 computer is receiving data from a PC and it sets the PC's CTS line false to tell it to stop sending data, the PC may send a few more bytes before actually stopping. Those bytes could be lost if you don't have a buffer.

The buffer might be compared to a holding tank in a Culligan home drinking water system. The water doesn't get through the reverse-osmosis filter nearly as fast as you would want it to come out to fill a cup. It's a waste of time to stand there for several minutes waiting for a dribble of water to fill the cup. On the other hand, the water can keep coming even when you're not needing it; so it makes sense to have the water coming through continually and going to the holding tank so it's there when you want it. When the tank is full (or nearly full), the water stops coming in.

So how do you implement this in software?

Making the size of the buffer or "holding tank" to be 256 bytes is convenient because we can use X or Y as an index into the buffer. When it is incremented from 255 to 256, we get 0 in the low 8 bits and let the high byte go. It makes for the easiest automatic wrap-around. 256 bytes works out to be a practical buffer size for common use anyway.

The buffer will act kind of like a continuous loop of recording tape. As data is written, the write pointer is incremented, always staying ahead of the read pointer. If the write pointer were to get a full circle ahead of the read pointer and catch up to it again, data that has not been read yet would get overwritten and be lost. In other words, we'd get a buffer-overflow condition. As the main program reads the data, it goes incrementing the read pointer. But if it catches up to the write pointer, it must stop or it will re-read old data as if it were new. As long as the write and read processes stay within the limits of the buffer, there is freedom for either end to go in bursts or have pauses that are not synchronized with those at the other end.




Figure 6. Receive buffer, with 35 (\$23) bytes in buffer. It would take 220 (\$DC) more bytes to fill the buffer.

So first we set up the buffer.

```

BUFFER: .DFS    $100    ; Reserve 256 bytes of RAM for the buffer ring itself.
RD_PTR: .DFS     1      ; Reserve one byte of RAM for the read pointer
WR_PTR: .DFS     1      ; and one for the write pointer.

```

Now we'll have a short routine to initialize the buffer. It doesn't normally matter where you start, since it will act as a continuous loop. We'll just set the read pointer to be the same as the write pointer. Each pointer points to the next byte to be written or read. When they're the same, the main routine knows there's nothing to read. If the buffer is full, the write pointer is one byte behind the read pointer. If it actually caught up, there would be no way of knowing if there were 0 bytes to read or 256; so the number of bytes left to read has to be in the range of 0 to 255, not 256.

```

INIT_BUF: LDA    RD_PTR    ; Setting WR_PTR equal to RD_PTR initializes
                STA    WR_PTR ; the buffer, showing it to be empty.
                RTS

```

```
;-----
```

Now a routine to write a byte to the buffer, one to read a byte from the buffer, and one to tell how many unread bytes are in the buffer:

```
WR_BUF:  LDX  WR_PTR    ; Start with A containing the byte to put in the buffer.
          STA  BUFFER,X  ; Get the pointer value and store the data where it says,
          INC  WR_PTR    ; then increment the pointer for the next write.
          RTS

;-----
RD_BUF:  LDX  RD_PTR    ; Ends with A containing the byte just read from buffer.
          LDA  BUFFER,X  ; Get the pointer value and read the data it points to.
          INC  RD_PTR    ; Then increment the pointer for the next read.
          RTS

;-----
BUF_DIF: LDA  WR_PTR    ; Find difference between number of bytes written
          SEC           ; and how many read.
          SBC  RD_PTR    ; Ends with A showing the number of bytes left to read.
          RTS

;-----
```

These will be used by the routine below. This routine is named SERVICE_ACIA, as called by the short routine "ISR" shown much earlier in this discussion. SERVICE_ACIA is written here with the assumption that its getting called necessarily means no other possible interrupt sources besides the 6551 ACIA produced the interrupt. If this is not the case, minor modification will be needed.

For this discussion, we will also assume for the sake of simplicity that we're not transmitting, and that DCD and DSR are always "true," meaning they won't be causing any interrupts.

```
SERVICE_ACIA:
    PHA
    PHX

    LDA  ACIA_STAT    ; Get the contents of the ACIA's status register.
    BPL  end          ; If ACIA didn't call, just exit.
    AND  #00000111B   ; Check for error conditions by ANDing with 7.
    BNE  REPORT_ERR   ; If there was any error condx, go report it.

    LDA  ACIA_DATA    ; Get the data from the ACIA
    JSR  WR_BUF       ; and store it in the buffer.

    JSR  BUF_DIF      ; Now see how full the buffer is.
    CMP  #$F0         ; If it has less than 240 bytes unread,
    BCC  end          ; just exit the ISR here.

    LDA  #1           ; Else, tell the other end to stop sending data before
    STA  ACIA_COMM    ; the buffer overflows, by storing 1 in the ACIA's
                      ; command register. (See text.)
end:    PLX
        PLA           ; Restore X and A, and return from interrupt.
        RTI

;-----
```

I didn't include a listing for REPORT_ERR above since how you handle it will depend heavily on your own system and how you want to handle it, and because it gets away from the subject of interrupts.

Storing a 1 in the ACIA's command register in the fourth-to-last line above sets the RTS line false, telling the other computer sending the data to stop. We stopped here with 16 bytes left to go so that if it sends out a few more bytes after we've told it to stop, we still don't overflow the buffer.

Now that we made it "turn off the faucet" so to speak, your routine that reads the data will have to see how full the buffer is and tell the sending end to start sending data again when you get below a certain point. After you've read a byte from the buffer using RD_BUF above, do something like

```
JSR  BUF_DIF    ; How many bytes are left to read?
CMP  #$E0       ; Is it at least 224?
BCS  buf_full   ; If so, leave the sending end turned off.
LDA  #00001001B ; Else, tell the sending end that it's ok to start
STA  ACIA_COMM  ; sending data again, by setting its CTS line true.
buf_full:       ; Continue...
```

Let's see-- have we forgotten anything? Oh yes-- I guess it would help to have a routine to set up the ACIA.

```
SETUP_ACIA:
    STZ  ACIA_STAT ; Reset ACIA by storing 0 in its status register.
    JSR  INIT_BUF  ; Initialize the software receive buffer.
```

```

LDA #00011110B ; Set for 1 stop bit, 8 data bits, 9600 bps by
STA ACIA_CTRL ; storing the number in the control register.

LDA #00001001B ; No parity or rcv echo, RTS true, receive IRQ but no
STA ACIA_COMM ; transmit IRQ, set DTR true. Store in command register.

RTS
;-----

```

If you wish, you can of course change the baud rate, the number of data and stop bits, choose odd or even parity, and/or make the link full-duplex as well, able to both transmit and receive data at the same time. The conditions above were chosen to keep things simple.

Ok, so what do you do with the data you read from the buffer? It could be that you're bringing in an Intel hex file for programming an EPROM or microcontroller, it could be you're loading a program, whether pre-compiled or -assembled or just source code, it could be you're bringing data in from another computer... you get the picture. How you handle the incoming data depends entirely on your application.

3.2 INTERRUPT-DRIVEN ARBITRARY WAVEFORM GENERATOR

Here is another example of interrupt use. An arbitrary waveform generator plays back a set of samples in memory over and over to produce a waveform you have already stored in that range of memory. Since you store whatever values you want, it is not limited to the typical function generators' sine, triangle, and square-wave outputs, or even white or pink noise.

We will use a VIA's T1 to generate the interrupts at an even rate. Each time it times out and interrupts, we will feed another 8-bit sample to the parallel D/A converter through port A (PA), then increment the address for the next sample to read from memory and feed to the D/A converter, see if the address has gone out of the intended range, and put it back to the beginning if so. Feeding the sample to the D/A converter *before* doing these other things assures us of minimal jitter; ie, that there is very little variation in the amount of time between when the interrupt hits and when sample value is fed to the converter.

T1 values for interrupt rates:

$$f(\text{IRQ}) = f(\phi 2) / (n+2)$$

$$n = f(\phi 2) / f(\text{IRQ}) - 2$$

TARGET	ACTUAL	5 MHz $\phi 2$ T1 Latch
12 kHz	11.9904 kHz	\$19F
16 kHz	15.9744 kHz	\$137
20 kHz	Exact	\$F8
24 kHz	24.0385 kHz	\$CE
28 kHz	28.0899 kHz	\$B0
36 kHz	35.9712 kHz	\$89
40 kHz	Exact	\$7B
48 kHz	48.0769 kHz	\$66

Zero-page variables:

```

CUR_ADR:      .DFS 2 ; The only STA() there is is in ZP, so use this addr.
                ; This will store the current address.

```

Not necessarily in zero page:

```

BEG_ADR:      .DFS 2 ; + These 2 specify the address range where the samples
END_ADR:      .DFS 2 ; + are stored, & get initialized in START_ARB below.
PLAY_T1_VAL:  .DFS 2 ; Value to store in VIA T1 counter for desired
                ; interrupt rate.

ARB_GEN_SAMPLER: ; (This is the ISR)
    PHA
    LDA VIA_IFR ; See if the sample timer (VIA's T1) timed out.
    BMI ags$ ; If so, go play another sample.
    PLA ; Otherwise, restore A and
    JMP _____ ; go test the next highest priority interrupt source.

ags$: LDA VIA_T1CL ; Clear the timer interrupt flag early on.

```

```

        LDA (CUR_ADR)    ; Get the byte to play back
        STA VIA_PA       ; and store it to the D/A convertor connected to PA.

        INC CUR_ADR      ; Increment the low byte of the current playback
        BNE ags1$        ; address. If that made it a 0,
        INC CUR_ADR+1    ; then you'll need to increment the high byte as well.

ags1$: LDA CUR_ADR      ; Now you need to see if you've done the last sample.
        CMP END_ADR     ; First compare low byte to that of ending address.
        BNE end$        ; If it's different, branch to end.
        LDA CUR_ADR+1    ; Then compare high byte
        CMP END_ADR+1    ; to that of ending address.
        BNE end$        ; If it's different, branch to end.

        LDA BEG_ADR      ; If you've done the last sample, set the current
        STA CUR_ADR      ; address so the ISR will start back at the
        LDA BEG_ADR+1    ; beginning of the sample table again the
        STA CUR_ADR+1    ; next time T1 times out and interrupts.

end$:   PLA              ; Restore value that was in A when the interrupt hit.
        RTI              ; Exit the ISR.
;-----

; After START_ARB below is called, the ISR will keep playing the loop of
; samples in the background until T1 interrupts are stopped. PLAY_T1_VAL and
; BEG_ADR and END_ADR should already be set to the desired values, although
; they can be changed after the sampling is started.

START_ARB:
        LDA #$FF        ; (Running this once starts the signal generation.)
        STA VIA_DDRA     ; To init parallel D/A operation, just set PA
                        ; to be all output bits.

        LDA VIA_ACR      ; Set T1 to free-run with IRQ every time-out,
        ORA #$C0         ; toggling VIA3PB7 (at half the interrupt rate)
        STA VIA_ACR      ; to check speed on oscilloscope if desired.
        LDA VIA_DDRB     ; (Data direction of PB7 should be set to output
        ORA #$80         ; though if you want to use PB7 to see the rate.)
        STA VIA_DDRB

        LDA PLAY_T1_VAL  ; Set the T1 time-out rate;
        STA VIA_T1CL     ; for example, $7B for 40kHz @ 5MHz φ2
        LDA PLAY_T1_VAL+1 ; as per table above.
        STA VIA_T1CH

        LDA #$C0
        STA VIA_IER      ; Enable T1 time-out IRQ.

        RTS
; -----

STOP_ARB:
        SEI              ; (This stops the signal generation started with START_ARB.)
        LDA #$40
        STA VIA_IER
        NOP              ; The NOPs here give some time for the IRQ pull-up
        NOP              ; resistor to clear an IRQ condition the could have
        CLI              ; coincided with the STA, before we enable interrupts
        RTS              ; again. This is to prevent "ghost" interrupts.
;-----

```

Note that if you have some kind of indirect interrupt vector in RAM, you will have to store ARB_GEN_SAMPLER's address there before the STA VIA_IER in START_ARB, and restore the original address there before the CLI in STOP_ARB. That original interrupt vector address should be stored as the operand of the JMP instruction in ARB_GEN_SAMPLER if any other interrupt sources may still be active.

3.3 USING THE 6522'S SERIAL PORT FOR SIMPLE ANALOG OUTPUT

Interestingly, if you can get away with only 9 output voltage levels (which is slightly better than 3-bit resolution), you can do a similar kind of thing with the VIA's serial port in mode 100, shift out free-running at T2 rate. In this mode, the shift register just continues to shift out the last byte you stored in it, over and over, until you give it another value. You only need the CB2 (serial data) line, and not the CB1 (serial clock) line. The idea here is that anywhere from 0 to 8 bits are set, giving you nine levels. Using something like:

00000000	for 0, gives .000V output;
00010000	for 1, gives .625V output;
00100010	for 2, gives 1.250V output;

01010010 for 3, gives 1.875V output;
10101010 for 4, gives 2.500V output;
10101101 for 5, gives 3.125V output;
11011101 for 6, gives 3.750V output;
11101111 for 7, gives 4.375V output; and
11111111 for 8, gives 5.000V output.

Of course, these voltages are ideal for the 65C22 (CMOS). The 6522 (NMOS) voltages will differ. Non-ideal conditions will not contribute significantly to distortion since the proportions are left intact, but the amplitude of the output signal may be reduced slightly.

Staggering the bits as much as possible as shown above facilitates filtering so you can get the smoothest possible analog output waveform with just a first-order RC output filter. This D/A method is adequate for generating DTMF (dual-tone multiple-frequency signals for telephone dialing), as well as other low-resolution signals. It will do much better at intelligible speech than many toys which use only two bits and slow sample rates. To make it easiest to filter, set the T2 latches for 0 so you get the maximum shift rate. T1 can still be used to control the rate at which the sample value is updated. Your RC output filter might be configured like that in Figure 7.




Figure 7. Simplest 6522 serial port D/A converter output filter, with -3dB point at 4kHz.

APPENDIX A. INTERRUPT PERFORMANCE

Figure 8 shows all the various times that should be considered when evaluating overall interrupt performance. (Some processors' sequence of events may vary from the diagram.) *Interrupt latency* is generally the amount of time it takes for the processor to begin the ISR after the interrupt line goes true. Sometimes this is measured in terms of the interrupt sequence only, which is 7 clocks for the 6502. The interrupt sequence is depicted in Figure 8(C). (The single-clock latency trick is still coming.)

Just as often however, interrupt latency includes more than just the sequence. It may also include the amount of time to finish up an already-started machine-language instruction before the interrupt sequence can begin, as depicted in Figure 8(A), and additionally on some processors, the time to do a separate interrupt-acknowledge step as depicted in Figure 8(B). The 6502's simpler method does not require this interrupt-acknowledge step or the associated extra hardware.

The longest-executing 6502 instructions (indexed read-modify-write instructions where the indexing causes a page boundary to be crossed) take 7 clocks. The average is about 4 clocks, depending on what you're doing. It is logical to assume that on the average, the interrupt will hit about in the middle of the instruction, making it necessary to complete about two more $\phi 2$ cycles before the interrupt sequence can begin. This gives an average total latency of about 9 clocks for the 6502.




Figure 8. Times to consider when evaluating interrupt performance.

The logic levels of the 6502's interrupt input pins are sampled on the falling edge of $\phi 2$. The falling edge of $\overline{\text{IRQ}}$ or $\overline{\text{NMI}}$ can be near the end of the last $\phi 2$ -high portion of the last cycle of a machine-language instruction, and the interrupt sequence will still begin before $\phi 2$ rises again. With instructions as long as 7 clocks, this means the total latency can be anywhere from slightly over 7 clocks (the time taken for the interrupt sequence itself) to slightly over 14 (the time for the interrupt sequence plus the longest-executing instruction). I say "slightly over" because the interrupt input must fall at least 10 to 200 nanoseconds (depending on the speed rating of the processor) before $\phi 2$ falls in order to be detected for that cycle. If the interrupt input falls too late, it will not be detected until the next cycle. If that next cycle starts another instruction, then that instruction will have to

be finished before the interrupt sequence will begin.

Another part of the processor's performance that is particular to interrupts is the return-from-interrupt instruction. In the case of many other processors, even this instruction depicted by Figure 8(E) in the diagram, whether called RETI, IRET, REIT, RETFIE, etc., by itself takes longer than A through E combined on a 16 or 20MHz 65C02 with a minimal ISR (like just incrementing a single byte in RAM)!

It would be fair to point out that some of the processors in the comparison below take so long for their interrupt sequence because they push a large set of registers onto the stack. The 6502 only has A, X, and Y that will sometimes need to be saved on the stack, but it's probably not very common that you would need even all three of those in your ISR; so always pushing all of them would not be the best use of time.

Some other processors have an entire second set of duplicate registers so that in a single instruction you can switch which set you're using, instead of taking the time to push them all onto the stack and restore them at the end of the ISR. While this may have its use, it would be similar to a one-cell hardware stack, so it is not re-entrant and its use would definitely have its limitations. I don't have personal experience with these, but have read that problems like this have indeed kept the supposed advantages from materializing.

In my own experience, it seems to me that the 6502 would, more often than not, be more efficient if it pushed the accumulator's contents onto the stack in the interrupt sequence. But I'm sure the designers, intent on maximizing overall interrupt performance instead of treating interrupts almost as an afterthought, weighed all the pros and cons of various ideas before deciding to leave the accumulator stack storage to the PHA and PLA instructions.

I have tried to get information on many other 8- and 16-bit processors to compare to. It has been difficult to get complete information on all of these, but here are bits and pieces on some for comparison. Only the Cygnal microcontroller gives an off-the-shelf 65C02 running at maximum speed any serious competition.

The single capital letters in parentheses below indicate, to the best of my understanding, what part of the interrupt time is being measured, in terms of Figure 8. Some of the information came from ST Microelectronics applications note AN910, at www.st.com/stonline/books/pdf/docs/5039.pdf. Those are noted by "ST" beside them. Other figures come from data books, either on my shelves or online. Frequencies shown in MHz are the maximum speeds as far as I have information. cy=clocks, unless otherwise noted. Avg=average.

1802	RCA	? Most instructions take 16 clocks (6.4µs), some, 24 (9.6µs). 2.5MHz @ 5V.
8080	Intel	? (still waiting for information)
8088	Intel	10 bus cycles or 40 clocks(?) (B)+(C) (still waiting for further information)
8086	Intel	WDC says 182 clocks max total latency. * * * * (still waiting for information)
Z8	Zilog	IRET (E) takes 16 execution cycles. I don't know how many clock cycles per execution cycle. 8MHz?
Z80	Zilog	11-19 clocks (B)+(C) depending on mode, or 2.75-4.75µs @ 4MHz. RETI (E) is 14 clocks, or 3.5µs @ 4MHz.
Z8000	Zilog	IRET (E) takes 13 cycles in non-segmented mode, and 16 in segmented mode. I don't know if that's instruction cycles or clock cycles.
8048	Intel	(?) return (E) is 2.7µs @ 11MHz
8051	Dallas	1.5µs @ 33MHz (52 clocks) latency
8051	Intel	1.8µs (C) min @ 20MHz. 5.4µs (A)+(C) max total latency @ 20MHz. (3-9µs @ 12MHz.) Interrupt sequence (C) and return (E) take 4.6µs @ 20MHz ST
80C51XA	Philips	2.25µs for interrupt+return (C)+(E) @ 20MHz. ST Instructions 2-24 cy, or 0.1-1.2µs. Avg 5-6 cy, or around 0.27µs.
KS88	Samsung	3µs for interrupt+return (C)+(E) @ 8MHz ST Instructions 6-28 cy, or 0.75-2.5µs. Avg 11 cy, or 1.38µs.
78K0	NEC	4.3µs for interrupt+return (C)+(E) @ 10MHz ST Instructions 4-50 cy, or 0.4-5.0µs. Avg 15 cy, or 1.5µs.
COP8	National	70 clocks (7 instruction cycles). RETI (E) is 50 clocks (5 instruction cycles). (7µs & 5µs @ 10MHz)

μPD78C05	NEC	RETI (E) takes 13 or 15 clocks (2.08 or 2.4μs at 6.25MHz)
μPD70008/A	NEC	sequence (C) takes 13 or 19 clocks. Return (E) takes 14 clocks. Instructions take 4-23 clocks each. 6MHz in '87 book.
V20	NEC	RETI (E) takes 39 clocks or 3.9μs @ 10MHz in '87 book. Instruction set is a superset of that of 8086/8088.
V25	NEC	? (still waiting for information)
68000	Motorola	46 clocks or 2.875μs minimum @ 16MHz (B)+(C)?. Has a very complex interrupt system.
6800	Motorola	(C)=13 clocks, including pushing the index register and both accumulators. RTI (E) takes 10 clocks. 2MHz.
6809	Motorola	(C)=19 clocks. Stacks all registers. RTI (E) 15 clocks. 2MHz (8MHz/4). FIRQ-RTI take 10 & 6 clocks, & work more like 6502 IRQ-RTI.
68HC05	Motorola	16 clocks typ (8μs @ 2MHz)
68HC08	Motorola	Instructions 1-9 cy, or 0.125-1.125μs. Avg 4-5 cy, or around 0.55μs.
68HC11	Motorola	(C)=14 clocks. RTI (E)=12 clocks. Total for interrupt+return=8.75μs @ 4MHz (16MHz/4). ST Instructions 2-41 cy, or 0.5-10.25μs. Avg 6-7 cy, or around 1.6μs.
68HC12	Motorola	2.63μs for interrupt+return (C)+(E) @ 8MHz. ST Instructions 1-13 cy, or 0.125-1.625μs. Avg 3-4 cy, or 0.45μs.
68HC16	Motorola	2.25μs for interrupt+return (C)+(E) @ 16MHz. ST Instructions 2-38 cy, or 0.125-2.375μs. Avg 6-7 cy, or around 0.4μs.
PIC16	Microchip	(C)=8 clocks (2 instruction cycles), and RETFIE (E) is also 8 clocks; but this doesn't even include saving and restoring the status register. That's an extra, rather mickey-mouse operation. 20MHz Most instructions 4 cy, or 0.2μs.
TMS370	Texas Instruments	15 cycles (3μs) min (C), 78 (15.6μs) max (A)+(C), and a cycle is 4 clocks (200ns min)! 20MHz. RTI (E) is 12 cy (48 clocks or 2.4μs).
TMS7000	Texas Instruments	(C)=19 cycles min (17 if from idle status) 5MHz, 400ns min cycle time (IOW, interrupt sequence is 7.6μs min, 6.8 from idle.) RETI (E) is 9 cycles, or 3.6μs @ 5MHz.
ST6	ST Microelectronics	78 clocks min, or 9.75μs @ 8MHz to fetch interrupt vector. More to reach first ISR instruction. RETI is 26 clocks, or 3.25μs.
ST7	ST Microelectronics	3μs for interrupt+return @ 8MHz. ST Instructions 2-12 cy, or 0.25-1.5μs. Avg 4-5 cy, or around 0.55μs.
ST9	ST Microelectronics	External IRQ best case: 1.08μs @ 24MHz. NMI best case: 0.92μs. internal interrupts best case: 1.04μs. 2.25μs @ 24MHz for interrupt and return ST Instructions 6-38 instruction cy, or 0.5-3.67μs. Avg 17 cy, or around 1.4μs.
ST9+	ST Microelectronics	1.84μs @ 25MHz for interrupt and return, ST Instructions 2-26 instruction cy, or 0.16-1.04μs. Avg 11 cy, or around 0.9μs
H8/300	Hitachi	8/16-bit: 2.1μs @ 10MHz for interrupt and return ST Instructions 2-24 cy, or 0.2-3.4μs. Avg 5-6 cy, or around 0.55μs.
M16C, M30218	Mitsubishi/Renesas	18 cy min (C), or 1.125μs @ 16MHz w/ 16-bit data bus. 50 cy max (A)+(C). REIT is 6 cy, or 0.375μs. Dual register sets like the Z80. Max instruction length 30 cy.
CIP-51	Silicon Labs Cygnal	μC p/n C8051F2xx total latency 5-18 cy or 0.2-0.72μs @ 25MHz. RETI takes 5 cy, or 0.2μs. This is the only one I have data on here that gives the 6502 below any competition.
65C02	WDC	Normal latency (C) 7 clocks (0.35μs) min, 14 clocks (0.7μs) max (A)+(C). RTI 6 cy (0.3μs). 20MHz. Instructions 2-7 cy, or 0.1-0.35μs. Avg 4 cy, or 0.2μs. Special case: IRQ from WAI instruction with interrupt-disable bit I set: no more than 1



Last Updated Mar 17, 2014, outdated cartoons notwithstanding!